


Lecture 1

Introduction to Distributed Systems



Presentation Outline

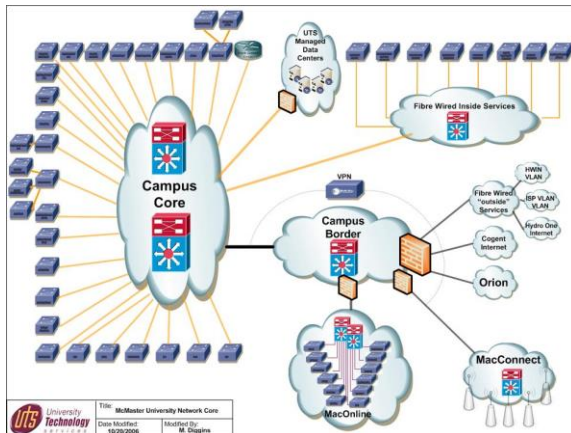

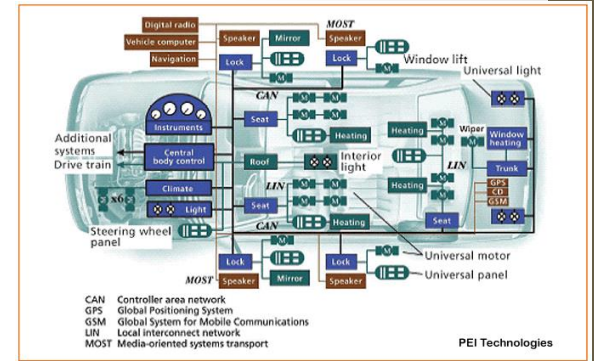
- Introduction
- Defining Distributed Systems
- Characteristics of Distributed Systems
- Example Distributed Systems
- Challenges of Distributed Systems
- Summary

Aims of this module

- Introduce the features of Distributed Systems that impact system designers and implementers
- Introduce the main concepts and techniques that have been developed to help in the tasks of designing and implementing Distributed Systems

Introduction

- Networks of computers are everywhere!
 - Mobile phone networks
 - Corporate networks
 - Factory networks
 - Campus networks
 - Home networks
 - In-car networks
 - On board networks in planes and trains




Defining Distributed Systems

- “A system in which hardware or software components located at *networked* computers communicate and coordinate their actions only by *message passing*.” [Coulouris]
- “A distributed system is a collection of *independent* computers *that appear* to the users of the system as a single computer.” [Tanenbaum]


Leslie Lamport's Definition

- *"A distributed system is one on which I **cannot** get any work done because some machine I have never heard of has crashed."*
- Leslie Lamport – a famous researcher on timing, message ordering, and clock synchronization in distributed systems.



Networks vs. Distributed Systems

- Networks: A media for interconnecting local and wide area computers and exchange messages based on protocols. Network entities are visible and they are explicitly addressed (IP address).
- Distributed System: existence of multiple autonomous computers is transparent
- However,
 - many problems (e.g., openness, reliability) in common, but at different levels.
 - Networks focuses on packets, routing, etc., whereas distributed systems focus on applications.
 - Every distributed system relies on services provided by a computer network.



Reasons for having Distributed Systems

- Functional Separation:
 - Existence of computers with different capabilities and purposes:
 - Clients and Servers
 - Data collection and data processing
- Inherent distribution:
 - Information:
 - Different information is created and maintained by different people (e.g., Web pages)
 - People
 - Computer supported collaborative work (virtual teams, engineering, virtual surgery)
 - Retail store and inventory systems for supermarket chains)

Reasons for having Distributed Systems

- Power imbalance and load variation:
 - Distribute computational load among different computers.
- Reliability:
 - Long term preservation and data backup (replication) at different locations.
- Economies:
 - Sharing resources to reduce costs and maximize utilization (e.g. network printer)
 - Building a supercomputer out of a network of computers.

Characteristics of Distributed Systems


- Concurrency
 - Carry out tasks independently and parallelly
 - Tasks coordinate their actions by exchanging messages
- Communication via message passing
 - No shared memory
- Resource sharing
 - Printer, database, other services
- No global state
 - No single process can have knowledge of the current global state of the system

Characteristics of Distributed Systems

- Heterogeneity – Different devices operating together
- Independent and distributed failures
- No global clock
 - Only limited precision for processes to synchronize their clocks

Differentiation with parallel systems

- Multiprocessor/Multicore systems
 - Shared memory
 - Bus-based interconnection network
 - E.g. SMPs (symmetric multiprocessors) with two or more CPUs, GPUs
- Multicomputer systems / Clusters
 - No shared memory
 - Homogeneous in hard- and software
 - Massively Parallel Processors (MPP)
 - Tightly coupled high-speed network
 - PC/Workstation clusters
 - High-speed networks/switches-based connection.




Differentiation with parallel systems is blurring

- Extensibility of clusters leads to heterogeneity
 - Adding additional nodes as requirements grow
- Leading to the rapid convergence of various concepts of parallel and distributed systems

Examples of Distributed Systems

- They (DS) are based on familiar and widely used computer networks:
 - Internet
 - Intranets, and
 - Wireless networks
- Example DS:
 - Web (and many of its applications like Facebook)
 - Data Centers and Clouds
 - Mobile applications
 - Wide area storage systems
 - Banking Systems



Challenges with Distributed Systems

- Heterogeneity
 - Heterogeneous components must be able to interoperate
- Distribution transparency
 - Distribution should be hidden from the user as much as possible
- Fault tolerance
 - Failure of a component (partial failure) should not result in failure of the whole system
- Scalability
 - System should work efficiently with an increasing number of users
 - System performance should increase with inclusion of additional resources

Challenges with Distributed Systems

- Concurrency
 - Shared access to resources must be possible
- Openness
 - Interfaces should be publicly available to ease inclusion of new components
- Security
 - The system should only be used in the way intended

Heterogeneity

- Heterogeneous components must be able to interoperate across different:
 - Operating systems
 - Hardware architectures
 - Communication architectures
 - Programming languages
 - Software interfaces
 - Security measures
 - Information representation



Mac OS



Distribution Transparency

- To hide from the user and the application programmer the separation/distribution of components, so that the system is perceived as a whole rather than a collection of independent components.
- ISO Reference Model for Open Distributed Processing (ODP) identifies the following forms of transparencies:
 - **Access transparency**
 - Access to local or remote resources is identical
 - E.g. Network File System / **Dropbox**
 - **Location transparency**
 - Access without knowledge of location
 - E.g. separation of domain name from machine address.



Distribution Transparency II


- Failure transparency
 - Tasks can be completed despite failures
 - E.g. message retransmission, failure of a Web server node should not bring down the website
- Replication transparency
 - Access to replicated resources as if there was just one. And provide enhanced reliability and performance without knowledge of the replicas by users or application programmers.
- Migration (mobility/relocation) transparency
 - Allow the movement of resources and clients within a system without affecting the operation of users or applications.
 - E.g. switching from one name server to another at runtime; migration of an agent/process from one node to another.

Distribution Transparency III

- **Concurrency transparency**
 - A process should not notice that there are other sharing the same resources
- **Performance transparency:**
 - Allows the system to be reconfigured to improve performance as loads vary
 - E.g., dynamic addition/deletion of components, switching from linear structures to hierarchical structures when the number of users increases
- **Scaling transparency:**
 - Allows the system and applications to expand in scale without changes in the system structure or the application algorithms.
- **Application level transparencies:**
 - Persistence transparency
 - Masks the deactivation and reactivation of an object
 - Transaction transparency
 - Hides the coordination required to satisfy the transactional properties of operations

Fault Tolerance

- Failure: an offered service no longer complies with its specification
- Fault: cause of a failure (e.g. crash of a component)
- Fault tolerance: no failure despite faults



Fault Tolerance Mechanisms

- Fault detection
 - Checksums, heartbeat, ...
- Fault masking
 - Retransmission of corrupted messages, redundancy, ...
- Fault toleration
 - Exception handling, timeouts,...
- Fault recovery
 - Rollback mechanisms,...

Scalability

- System should work efficiently at many different scales, ranging from a small Intranet to the Internet
- Remains effective when there is a significant increase in the number of resources and the number of users
- Challenges of designing scalable distributed systems:
 - Cost of physical resources
 - Performance Loss
 - Preventing software resources running out:
 - Numbers used to represent Internet addresses (32 bit->64bit)
 - Y2K-like problems
 - Avoiding performance bottlenecks:
 - Use of decentralized algorithms (centralized DNS to decentralized)

Concurrency

- Provide and manage concurrent access to shared resources:
 - Fair scheduling
 - Preserve dependencies (e.g. distributed transactions)
 - Avoid deadlocks
 - Preserve integrity of the system

Java Concurrency



Openness and Interoperability

- Open system:
"... a system that implements sufficient **open specifications** for interfaces, services, and supporting formats to enable properly engineered applications software to be ported across a wide range of systems with minimal changes, to interoperate with other applications on local and remote systems, and to interact with users in a style which facilitates user portability" (Guide to the POSIX Open Systems Environment, IEEE POSIX 1003.0)

Openness and Interoperability


- Open message formats: e.g. XML
- Open communication protocols: e.g. HTTP, HTTPS
- Open spec/standard developers - communities:
 - ANSI, IETF, W3C, ISO, IEEE, OMG, Trade associations,...

Security I

- Resources are accessible to authorized users and used in the way they are intended
- Confidentiality
 - Protection against disclosure to unauthorized individual information
 - E.g. ACLs (access control lists) to provide authorized access to information
- Integrity
 - Protection against alteration or corruption
 - E.g. changing the account number or amount value in a money order

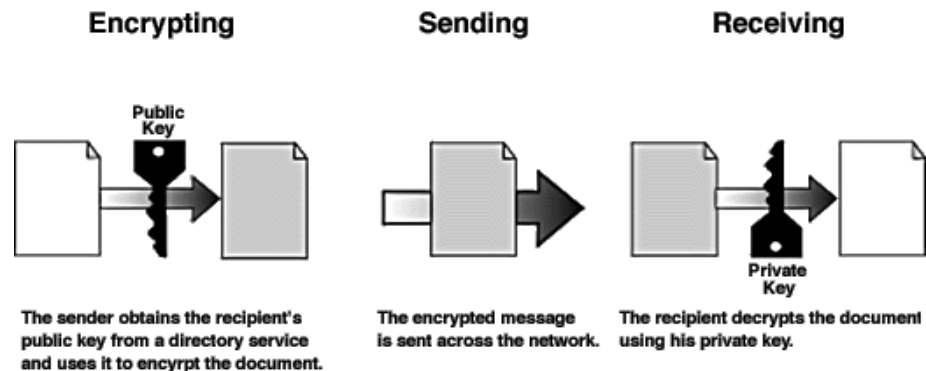
Security II

- Availability
 - Protection against interference targeting access to the resources.
 - E.g. denial of service (DoS, DDoS) attacks
- Non-repudiation
 - Proof of sending / receiving an information
 - E.g. digital signature



Security Mechanisms

- Encryption
 - E.g. Blowfish, RSA
- Authentication
 - E.g. password, public key authentication
- Authorization
 - E.g. access control lists



Business Example and Challenges

- Web/Mobile app to search and purchase online courses
 - Customers can connect their computer to your server (locally hosted or cloud hosted):
 - Browse your courses
 - Purchas courses
 - ...

Business Example – Challenges

I

- What if
 - Your customers use different devices? (Dell laptop, Android device ...)
 - Your customers use different OSs? (Android, iOS, Ubuntu...)
 - a different way of representing data? (Text, Binary,...)
 - **Heterogeneity**
- Or
 - You want to move your business and computers to China (because of the lower costs)?
 - Your client moves to a different country(more likely)?
 - **Distribution transparency**

Business Example – Challenges


II

- What if
 - Two customers want to order the same item at the same time?
 - **Concurrency**
- Or
 - The database with your inventory information crashes?
 - Your customer's computer crashes in the middle of an order?
 - **Fault tolerance**

Business Example – Challenges

III

- What if
 - Someone tries to break into your system to alter data?
 - ... sniffs for information?
 - ... someone says they have enrolled to the course but they haven't?
 - **Security**
- Or
 - You are so successful that millions of people are using your app at the same time.
 - **Scalability**



Business Example – Challenges

IV

- When building the system...
 - Do you want to write the whole software on your own (network, database,...)?
 - What about updates, new technologies?
 - Adding a web client later on?
 - Will your system need to communicate with existing systems (e.g. payment gateways, SMS servers)
 - **Reuse** and **Openness** (Standards)



Impact of Distributed Systems

- New business models (e.g. Uber, Airbnb)
- Global financial markets
- Global labor markets
- E-government (decentralized administration)
- Ecommerce
- Driving force behind globalization
- Social/Cultural impact
- Media getting decentralized

Summary

- Distributed Systems are everywhere
- The Internet enables users throughout the world to access its services wherever they are located
- Resource sharing is the main motivating factor for constructing distributed systems
- Construction of DS produces many challenges:
 - Heterogeneity, Openness, Security, Scalability, Failure handling, Concurrency, and Transparency
- Distributed systems enable globalization:
 - Community (Virtual teams, organizations, social networks)
 - Science (e-Science)
 - Business (e-Business)